

# Yerram Varun

 y.varun@nyu.edu

 varun221.github.io

 LinkedIn

 Google Scholar

 GitHub

## EDUCATION

---

**New York University, Center for Data Science**

2025 - Present

*Ph.D. Data Science*

*Advisors: Prof. Eunsol Choi, Prof. He He*

**Indian Institute of Technology, Guwahati**

2019 - 2023

*Bachelor of Technology in Electronics & Communication*

*Minor in Computer Science & Engineering*

## RESEARCH EXPERIENCE

---

**Google DeepMind, India**

Aug 2023 - July 2025

*Pre-Doctoral Researcher*

*Advisors: Dr. Prateek Jain ↗, Dr. Praneeth Netrapalli ↗ & Dr. Karthikeyan Shanmugam ↗*

### Time Reversed LLMs

- Implemented pre-training & post-training pipelines with token reversal data processing to train Reversed-LLMs of LaMDA, PaLM2 and Gemini Family of models.
- Performed exhaustive experiments into various applications viz Reranking, Citations, Retrieval, Safety and Planning.
- Showed extensive improvements in Instruction following abilities, Citation attribution and Passage Retrieving capabilities of Gemini-1.0 Pro resulting in **Spotlight submission @ NeurIPS, 2024**.

### Efficient Inference: High Recall Estimation

- Worked on efficient inference in FFN and Softmax layers of LLM through **sparsity exploitation** techniques.
- Performed extensive ablation studies to design algorithms that **cheaply** predict top non-zero rows/columns of activations with **high recall** in distributed inference settings.
- We achieve **1.47x** per step decode speedup with minimal finetuning on 1 billion parameter models.

### Long Context Efforts : Gemini Family of Models

- Worked on approximate logit computation and clustering approaches to identify top-k keys in the attention layer.
- Designed approaches to retain quality and improve latency with existing methods in a distributed inference setting.
- Showed significant latency improvements on Gemini at different scales and settings.
- Actively contributing to the productionization efforts to improve current Long-context Gemini models.

**Indian Institute of Technology, Guwahati**

Jan 2022 - Mar 2022

*Collaborator: Prof. Vivek Gupta ↗*

### Trans-KBLSTM: Enhancing Tabular Natural Language Inference

- Developed custom architectures using **RoBERTa** and **BiLSTMs** to integrate external knowledge bases like **ConceptNet** and **Wordnet** into attention architectures.
- Improved SOTA Accuracy on Infotabs, a tabular Natural Language Inference Dataset by 4-5% on low-resource settings.
- Published and received **Best Paper Award** at DeeLIO Workshop@ACL 2022

## PUBLICATIONS

---

- **Time-Reversal Provides Unsupervised Feedback to LLMs.** ↗  
**Y Varun\***, R Madhavan\*, S Addepalli\*, A Suggala, K Shanmugam, P Jain.  
*Neural Information Processing Systems, 2024.*

 **Spotlight Submission.**

- **Does Safety Training of LLMs Generalize to Semantically Related Natural Prompts?** ↗  
S Addepalli, **Y Varun**, A Suggala, K Shanmugam, P Jain.  
*SafeGenAI Workshop @ Neural Information Processing Systems, 2024*  
*Under submission to ICLR, 2025*

- **HiRE: High Recall Approximate Top-k Estimation for Efficient LLM Inference** ↗  
Y Samaga B L\*, **Y Varun\***, C You, S Bhojanapalli, S Kumar, P Jain, P Netrapalli
- **Trans-KBLSTM: An External Knowledge Enhanced T-BiLSTM Model for Tabular Reasoning** ↗  
**Y Varun**, A Sharma, V Gupta  
*DeeLIO Workshop @ Association for Computational Linguistics, 2022*  
 **Best Paper Award.**

- **FlowVQA: Mapping Multimodal Logic in Visual Question Answering with Flowcharts** ↗  
S Singh\*, P Chaurasia\*, Y Varun, P Pandya, V Gupta, V Gupta, D Roth  
*Findings of Association for Computational Linguistics 2024*

## ENGINEERING EXPERIENCE

---

**India Machine learning Team, Amazon Science, Bangalore** ↗

May 2022 - July 2022

*Applied Scientist Intern*

- Worked on (i) Top-K feature-selection and (ii) Temporal learning based attention architectures to model customer demographics through Amazon Behavioral history.
- Improved the existing XGBoost production pipeline by 2-3% F1 score with **Time2Vec** and **Multi-head Attention** based algorithms and contributed to initial efforts to productionize.

**RethinkUX, India** ↗

Jan 2021 - Feb 2021

*DL Software Development Intern*

- Developed **text detection models** and **text recognition methods** to extract information from PDF documents.
- Implemented parsing-based pipelines to extract required customer attributes from parsed documents.
- Created a **Flask API** and used **Docker** to containerise the application for easy deployment over system.
- Created a complete **ML pipeline** from data preparation to deployment in production.

## ACHIEVEMENTS

---

- **Peer Bonus** ↗, For providing guidance on coding infrastructure & Model/Data workflows @ Google DeepMind, India
- **Peer Bonus** ↗, For providing guidance on dataset pipelines @ Google DeepMind, India
- **Best Paper Award DeeLIO Workshop, Association for Computational Linguistics 2022**

### Hackathons

- **Amazon Machine learning Challenge, 2021**, National Runner up among 10,000+ registered participants
  - **Inter IIT Tech Meet 10.0**, Secured Silver Medal in Bosch Model Extraction Module
  - **Inter IIT Tech Meet 9.0**, Secured Silver Medal in the Bosch Traffic Sign Recognition Module
- ### Kaggle ML Competitions
- **RSNA 2022 Cervical Spine Fracture Detection - Kaggle**, Finished 38th (Silver Region), Top 5%
  - **Mayo Strip AI Competition - Kaggle**, Finished 27th (Silver Region), Top 4%
  - **DFL : Bundesliga Data Shootout - Kaggle**, Finished 44th (Silver Region), Top 8%
  - **Sartorius Cell Instance Segmentation - Kaggle**, Finished 111th (Bronze Region), Top 8%
  - **Mechanisms of Action - Kaggle**, Finished 208th (Silver Region), Top 5%

## RELEVANT COURSEWORK

---

**Mathematics** - Linear Algebra, Multivariate Calculus, Probability & Random Processes(AA Grade)

**Computer Science** - Deep Learning (AA Grade), Data Structures & Algorithms, Machine Learning, Computer Vision

## TECHNICAL SKILLS

---

*Languages* (Python, C/C++, MATLAB, LaTeX, HTML/CSS, JS, Bash, SQL, MATLAB), *Libraries* (JAX, Pytorch, Tensorflow, Keras, FastAI, PySpark, Flask, Django), *Tools* (Git, HuggingFace, Bootstrap)

## EXTRACURRICULARS

---

- **Secretary, IITG.ai**, Responsible for managing events, workshops, competitions, and discussions at IITG.ai club, the AI Community of IIT Guwahati
- **Speaker at MLTechFest' 21** ↗, Conducted a Workshop on basics of transformers at MLTechFest'21 organized by Tensorflow User Group (Mysuru) & Google Developer's Group (Mysuru)
- **Computer Vision Workshop by IEEE SFIT & IEEE APSIT** ↗, Conducted an online workshop session on basics of Computer Vision and OpenCV
- **First Position in Robothon 2020 organized by Robotics Club, IITG**
- **First Position in Astronomy Quiz, Interhostel Technical Competition, Kriti 2020, IITG**